

WAYS TO PRODUCTIONIZE H2O

Tom Kraljevic
tomk@h2o.ai

H2O Open Tour 2016, New York City
July 19, 2016

WHERE TO FIND THE SLIDES

They're in GitHub...

[h2oai/h2o-meetups/tree/master/
2016_07_19_H2O_Open_Tour_NYC_Prod](https://github.com/h2oai/h2o-meetups/tree/master/2016_07_19_H2O_Open_Tour_NYC_Prod)

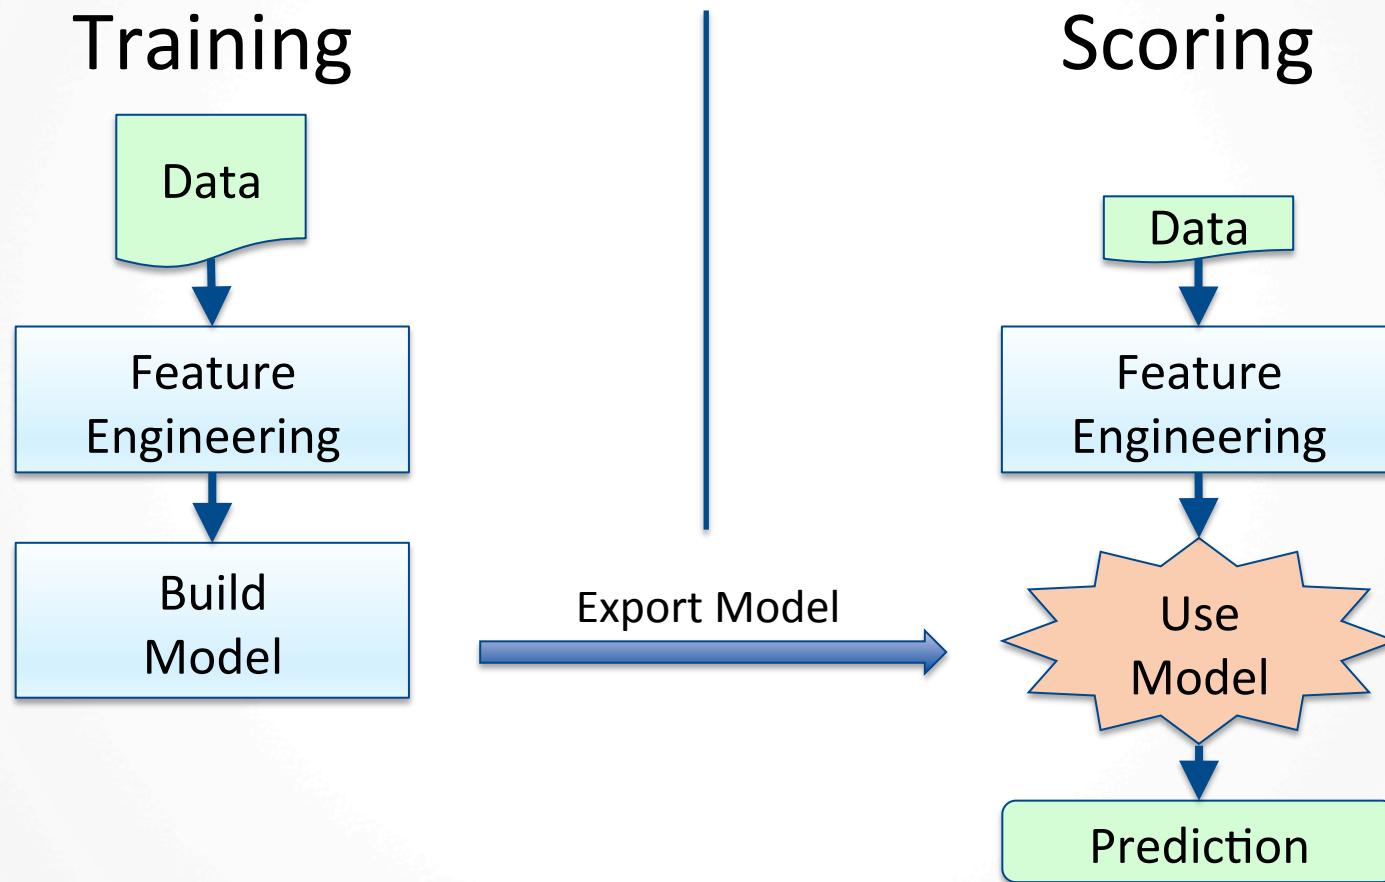
DEFINING “IN PRODUCTION”?

- Removal would impact value to the business
 - Dollars
 - Time

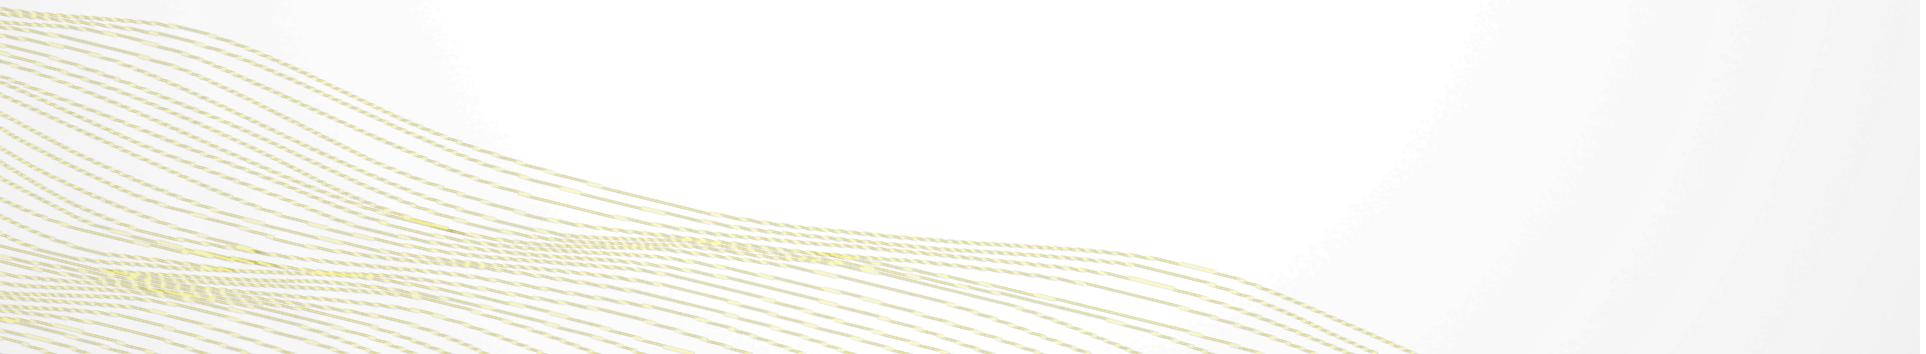
OUTLINE

- Motivation
- Productionizing Model Building
 - Training
- Productionizing Model Use
 - Scoring aka Predicting
- Q & A

TYPICAL DEPLOYMENT APPROACH



PRODUCTIONIZING MODEL BUILDING (TRAINING)



H2O TRAINING DIMENSIONS

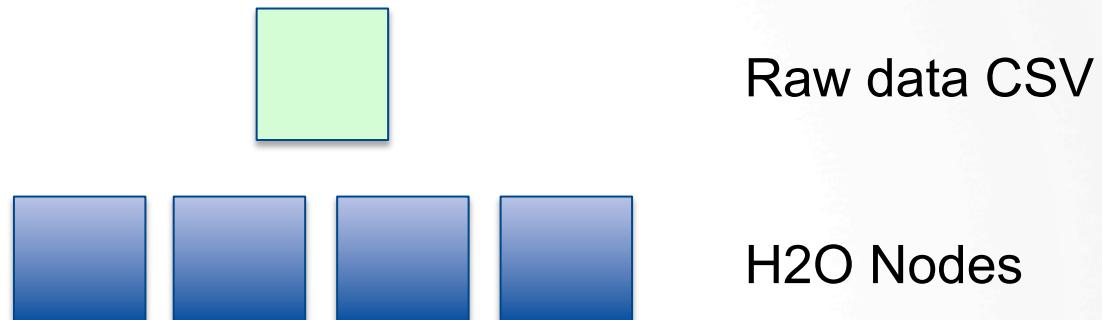
- Where does the data live?
 - Hadoop
 - HDFS, Hive
 - S3 / AWS
 - Database
 - Local drive (laptop, desktop, server)
- How big is the data?
 - Fits in memory of one host? (Or not)
 - “Too small”?
- How often should models be re-trained?
 - Daily? Weekly? Monthly? Quarterly?
- Where will model training happen?
 - Hadoop YARN job
 - EC2 instances
 - Bare metal cluster
 - A really big server
 - Laptop
- Resource requirements (how many nodes?)
 - CPUs
 - Memory
 - Time to build a model

H2O TRAINING SIZING BEST PRACTICE

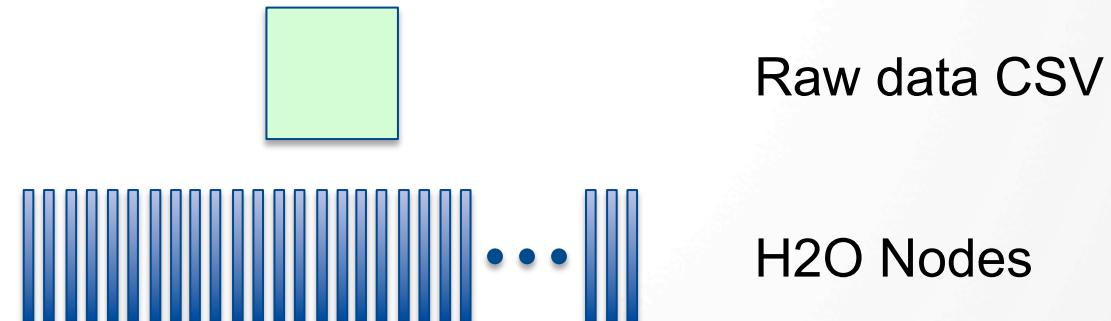
- For a novice getting started with a decent sized dataset:
 - Size total cluster memory to 4x on-disk data size (data for model training, not “total data”)
 - Use nodes of at least 10 GB, preferably larger
 - Use a small number of big nodes
 - e.g. 8 nodes x 30 GB each
 - For a YARN environment, you will have to adjust default settings to start big containers
 - **Definitely DO NOT use 200 nodes x 1 GB each**
 - Caution: Hadoop MapReduce veterans often start this way!
 - 1 GB is way too small for *any* workload!

H2O TRAINING SIZING BEST PRACTICE

Start
With This



For the
Experts

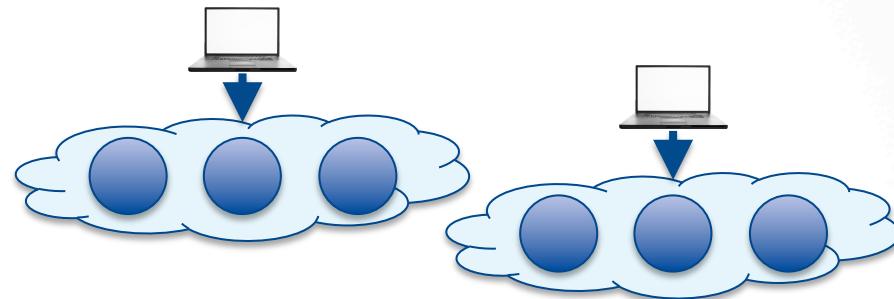


H2O USER PROVISIONING BEST PRACTICE

- Individual H2O jobs (aka application instances, aka clusters) are not multi-tenant
 - Think of YARN as providing multi-tenancy and H2O as a single job
 - H2O instance per interactive user, or per batch job
 - Steam product from H2O.ai simplifies management of H2O clusters

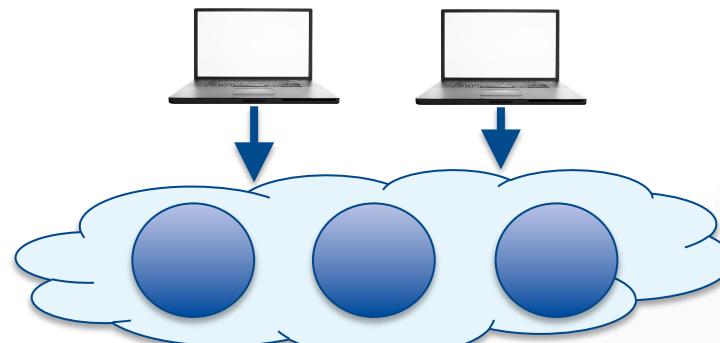
H2O USER PROVISIONING BEST PRACTICE

Recommended



1 User per
H2O Instance

Not
Recommended



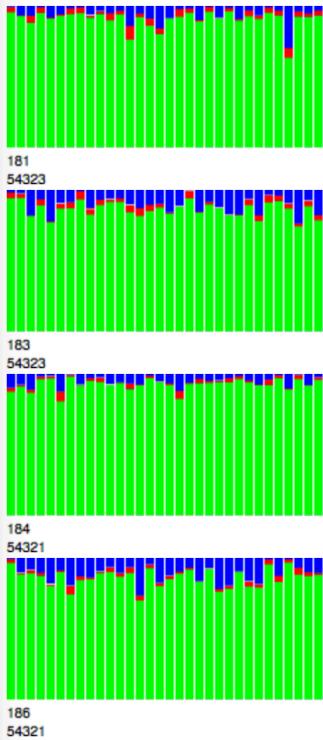
Shared Group
H2O Instance

H2O MONITORING BEST PRACTICE

- Monitor your H2O instance's resource utilization
 - Watch the Water Meter in the Flow Web UI to see:
 - If your hardware cluster is already being slammed before H2O even starts doing anything
 - If your H2O instance is using CPU when you expect it to
 - Look at H2O log files for detailed memory utilization information (advanced)

H2O MONITORING BEST PRACTICE

CPU Monitoring



Memory Monitoring (advanced)

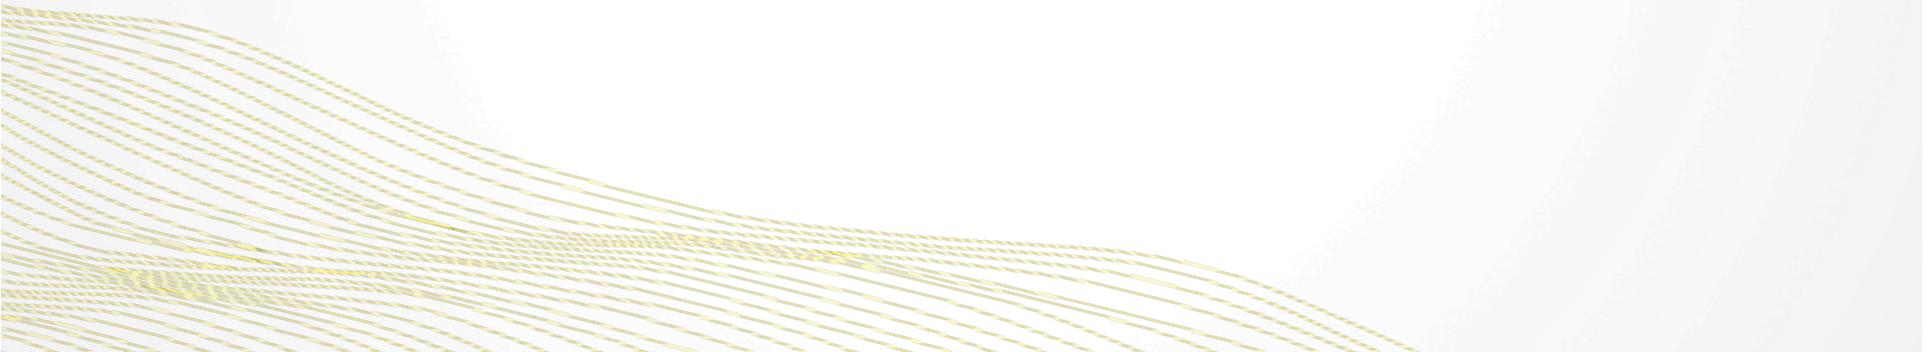
yarn logs -applicationId application_1457562501251_2398

```
07-18 20:18:03.364 172.16.2.184:54321 31423 main INFO: ---- H2O started ----
07-18 20:18:03.430 172.16.2.184:54321 31423 main INFO: Build git branch: rel-turchin
07-18 20:18:03.430 172.16.2.184:54321 31423 main INFO: Build git hash: c09d9da94e7c1eb7ef4573f9f5029e0ca501469
07-18 20:18:03.431 172.16.2.184:54321 31423 main INFO: Build git describe: jenkins-rel-turchin-3
07-18 20:18:03.431 172.16.2.184:54321 31423 main INFO: Build project version: 3.8.2.3
07-18 20:18:03.431 172.16.2.184:54321 31423 main INFO: Built on: '2016-04-25 13:07:53'
07-18 20:18:03.431 172.16.2.184:54321 31423 main INFO: Java availableProcessors: 32
07-18 20:18:03.431 172.16.2.184:54321 31423 main INFO: Java heap totalMemory: 39.45 GB
07-18 20:18:03.431 172.16.2.184:54321 31423 main INFO: Java heap maxMemory: 39.45 GB
07-18 20:18:03.431 172.16.2.184:54321 31423 main INFO: Java version: Java 1.7.0_67 (from Oracle Corporation)
07-18 20:18:03.431 172.16.2.184:54321 31423 main INFO: JVM launch parameters: [-XX:+PrintGCTimeStamps, -Djava.net.preferIPv4Stack=true, -Dhdp.version=2.2.6.3-1, -Xms40g, -Xmx40g, -XX:PermSize=256m, -verbose:gc, -XX:+PrintGCDetails, -XX:+PrintGCTimeStamps, -Dlog4j.defaultInitOverride=true, -Djava.io.tmpdir=/opt2/hdp/yarn/local/usercache/tomk/appcache/application_1457562501251_2398/container_e13_1457562501251_2398_01_000005/tmp, -Dlog4j.configurationFile=/opt2/hdp/yarn/log/application_1457562501251_2398/container_e13_1457562501251_2398_01_000005/conf/log4j.properties, -Dyarn.app.container.log.dir=/opt2/hdp/yarn/log/application_1457562501251_2398/container_e13_1457562501251_2398_01_000005, -Dyarn.app.container.log.filesize=0, -Dhadoop.root.logger=INFO,CLAA]
07-18 20:18:03.431 172.16.2.184:54321 31423 main INFO: OS version: 3.13.0-61-generic (amd64)
07-18 20:18:03.431 172.16.2.184:54321 31423 main INFO: Physical memory: 251.89 GB
07-18 20:18:03.431 172.16.2.184:54321 31423 main INFO: X-h2o-cluster-id: 1468898282076
07-18 20:18:03.431 172.16.2.184:54321 31423 main INFO: User name: 'tomk'
07-18 20:18:03.432 172.16.2.184:54321 31423 main INFO: Possible IP Address: br2 (br2), 172.16.2.184
07-18 20:18:03.432 172.16.2.184:54321 31423 main INFO: Possible IP Address: lo (lo), 127.0.0.1
07-18 20:18:03.432 172.16.2.184:54321 31423 main INFO: Internal communication used port: 54322
07-18 20:18:03.432 172.16.2.184:54321 31423 main INFO: Listening for HTTP and REST traffic on http://172.16.2.184:54321
07-18 20:18:03.432 172.16.2.184:54321 31423 FJ-1-15 INFO: Frame distribution summary:
07-18 20:29:58.542 172.16.2.184:54321 31423 FJ-1-15 INFO: Number of Rows: 29811120, Number of Chunks per Column: 131.000000, Number of Chunks: 4061.000000
07-18 20:29:58.542 172.16.2.184:54321 31423 FJ-1-15 INFO: Number of Rows: 29023380, Number of Chunks per Column: 127.000000, Number of Chunks: 3037.000000
07-18 20:29:58.542 172.16.2.184:54321 31423 FJ-1-15 INFO: Number of Rows: 28900760, Number of Chunks per Column: 127.000000, Number of Chunks: 3037.000000
07-18 20:29:58.542 172.16.2.184:54321 31423 FJ-1-15 INFO: Number of Rows: 29173814, Number of Chunks per Column: 128.000000, Number of Chunks: 3096.000000
07-18 20:29:58.542 172.16.2.184:54321 31423 FJ-1-15 INFO: mean 911.8 MB
07-18 20:29:58.542 172.16.2.184:54321 31423 FJ-1-15 INFO: min 883.7 MB
07-18 20:29:58.542 172.16.2.184:54321 31423 FJ-1-15 INFO: max 925.5 MB
07-18 20:29:58.542 172.16.2.184:54321 31423 FJ-1-15 INFO: stddev 16.5 MB
07-18 20:29:58.542 172.16.2.184:54321 31423 FJ-1-15 INFO: total 3.56 GB
07-18 20:29:58.542 172.16.2.184:54321 31423 FJ-1-15 INFO: 1.732051 53.693577
07-18 20:29:58.542 172.16.2.184:54321 31423 FJ-1-15 INFO: 512.000000 15872.000000
793.532: [GC [PSYoungGen: 2307135K->1553395K(3107328K)] 6982139K->6408290K(40390144K), 0.1289179 secs] [Times: user=0.19 sys=0.53, real=0.12 secs]
793.532: [GC [PSYoungGen: 2307135K->1553395K(3107328K)] 7862139K->7350065K(40390144K), 0.1289750 secs] [Times: user=0.61 sys=0.17, real=0.17 secs]
794.524: [GC [PSYoungGen: 3106455K->1221115K(3107328K)] 89031217K->8571153K(40390144K), 0.1658988 secs] [Times: user=0.92 sys=0.77, real=0.17 secs]
795.501: [GC [PSYoungGen: 2774868K->1496653K(3107328K)] 10125206K->9965740K(40390144K), 0.1611300 secs] [Times: user=2.81 sys=0.82, real=0.16 secs]
796.684: [GC [PSYoungGen: 3050216K->986551K(3107328K)] 11519302K->10546831K(40390144K), 0.1470070 secs] [Times: user=2.03 sys=1.27, real=0.15 secs]
797.983: [GC [PSYoungGen: 2540021K->781681K(3107328K)] 120899782K->11006879K(40390144K), 0.1154810 secs] [Times: user=1.38 sys=1.21, real=0.12 secs]
```

ADDITIONAL YARN BEST PRACTICES

- Limiting memory occurs “naturally”
- No CPU usage limit by default, but you can use:
 - YARN cgroups
 - H2O –nthreads startup option
- Use node labels to direct container placement
- Capacity and Fair schedulers are both fine
- Many people provision YARN queues to manage resource use for teams
- You *must* disable preemption for H2O jobs

PRODUCTIONIZING MODEL USE (SCORING AKA PREDICTING)



HIGH-LEVEL H2O DIMENSIONS

In-H2O

- Batch
- Many rows at once
 - think Frames of data
- R / Python / Scala
 - script driving H2O
- Script

POJO

- Real-time
- One row at a time
- Java / Scala
 - embedded or as a service
- Application

OTHER INDEPENDENT DIMENSIONS

Choice A

- On-premise
- Data Science
- Micro-service with REST API
- Model only
- Train and predict together
- Pre-calculated

Choice B

- Cloud
- Computer Science
- Embedded
- Model + Feature engineering
- Train and predict separately
- Calculate live

DATA ENVIRONMENTS

- Hive
- Databases
 - Oracle, MS SQL, SQLite, etc.
- Cloud databases
(Amazon RDS)
- Teradata
- Bare flat files
- Hadoop HDFS
- Amazon S3
- Storm spout
- Spark stream
- Apache Apex

MODEL ENVIRONMENTS

- Laptop / Desktop
- Hadoop
- Standalone cluster
- Steam
- Cloud (e.g. EC2)
- Cloud endpoint (e.g. AWS Lambda)
- Storm bolt
- Spark stream
- Apache Apex
- Hive UDF
- POJO called from R

KEY PARAMETERS

- Model training
 - Frequency
 - Cost (memory, CPU, time)
- Prediction SLA
 - Throughput
 - Latency
 - Availability
- Security

EXAMPLE DESIGN PATTERNS

In-H2O

- Driven from Flow/R/
Python/Scala
 - Often launching on Hadoop
- Embedding H2O (droplet)

POJO

- JUnit test
- Craigslist app (Sparkling Water)
- App-consumer-loan (Jetty servlet)
- App-malicious-domain (AWS lambda)
- POJO from R (jar file)
- Journals.ai desktop app (SQLite DB)
- Storm bolt
- Hive UDF
- Steam from H2O.ai (REST API)

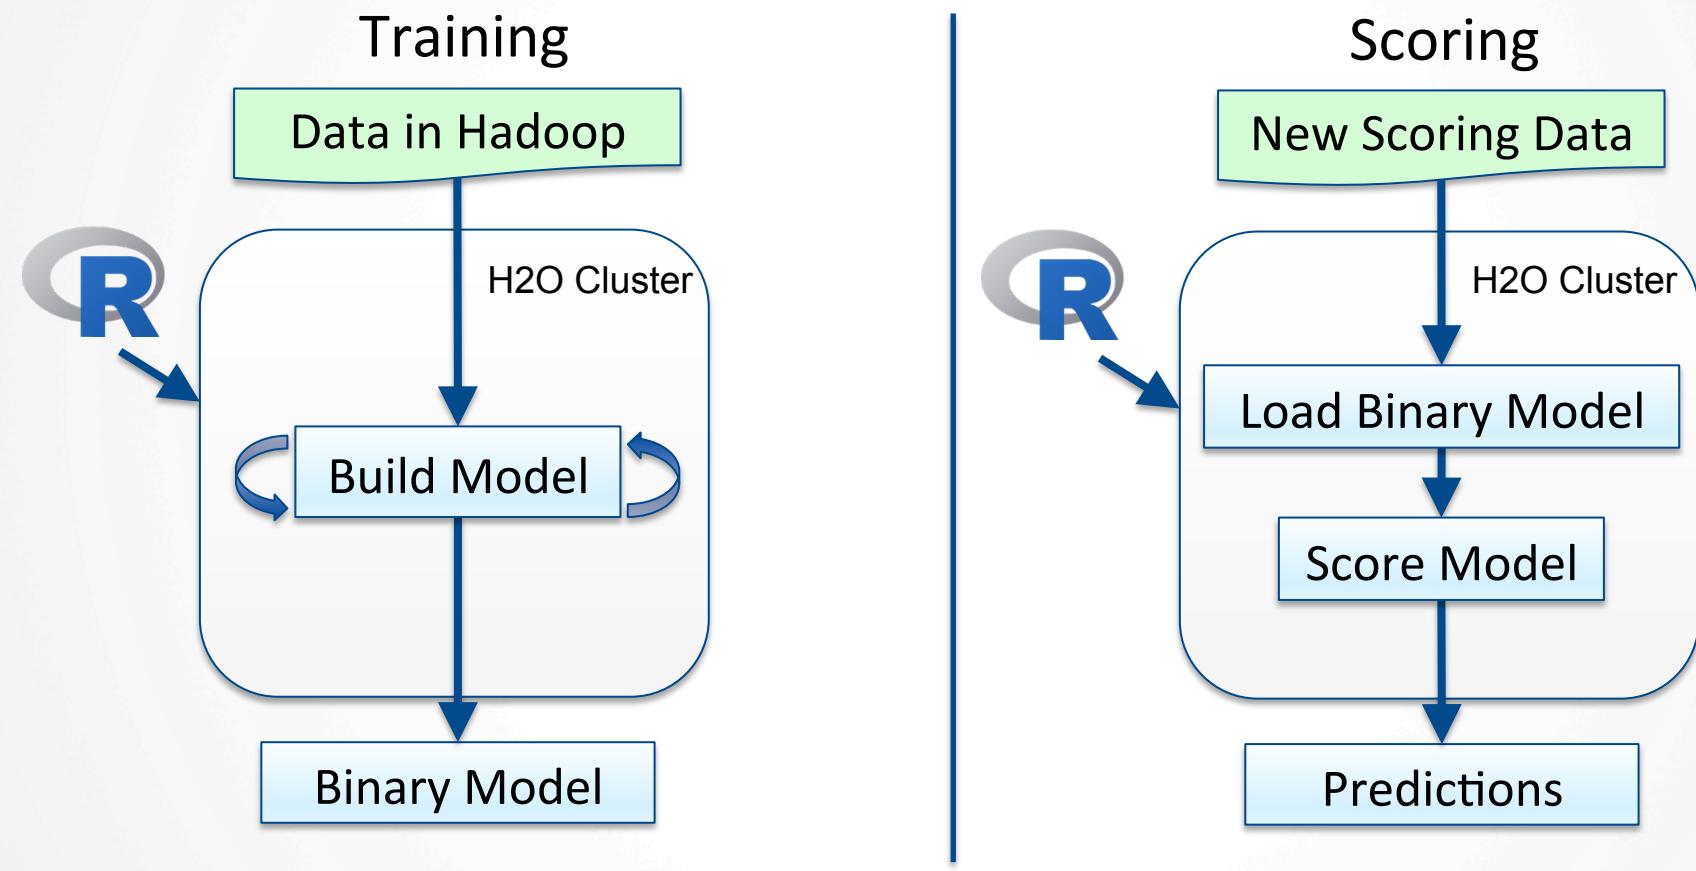
WALKTHROUGH FOR 5 DESIGN PATTERNS

- Batch In-H2O
- Steam
- AWS Lambda
- Hive UDF
- Pre-calculated DB

PURE BATCH H2O DESIGN PATTERN

- Training and scoring data lives in data lake
- Entire process driven by R
- Training occurs periodically
- Scoring occurs frequently

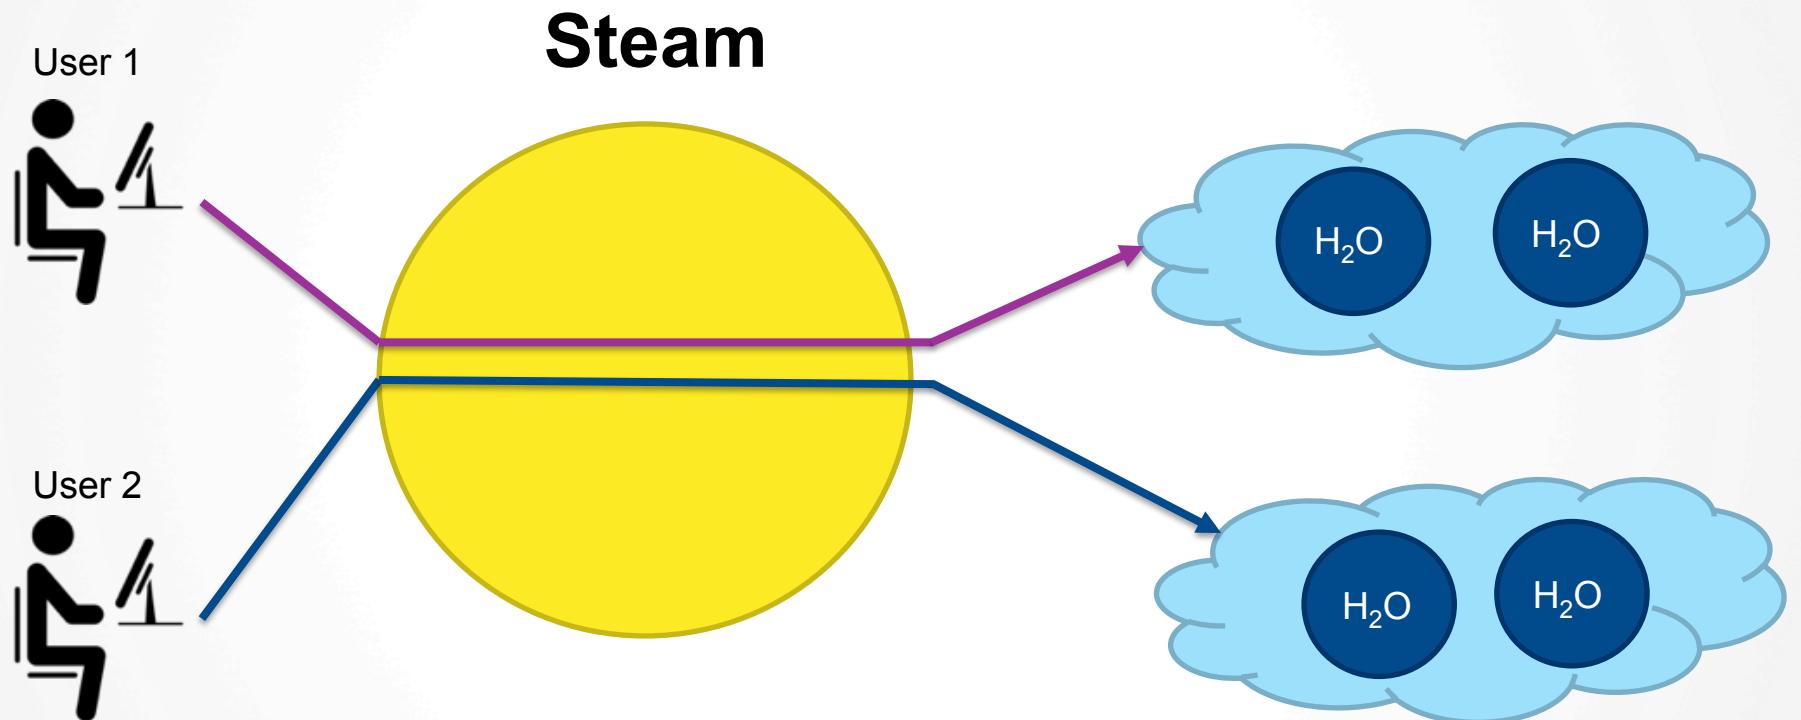
PURE BATCH H2O DESIGN PATTERN



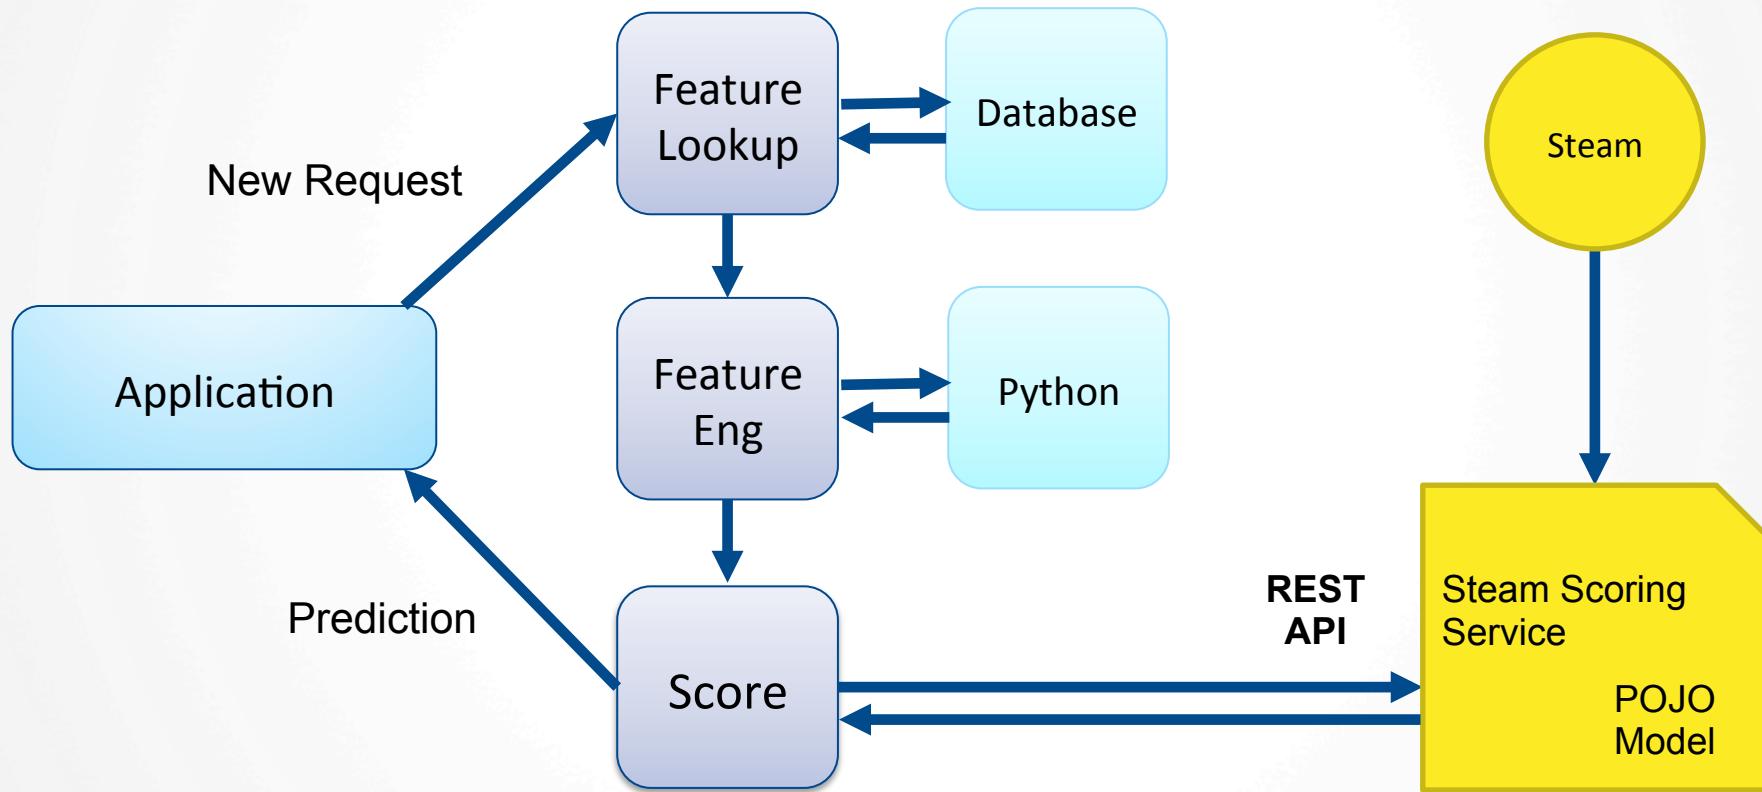
STEAM DESIGN PATTERN

- Training data lives in Hadoop
- POJO models are deployed as a service with a REST API

STEAM DESIGN PATTERN (TRAINING)



STEAM DESIGN PATTERN (SCORING)



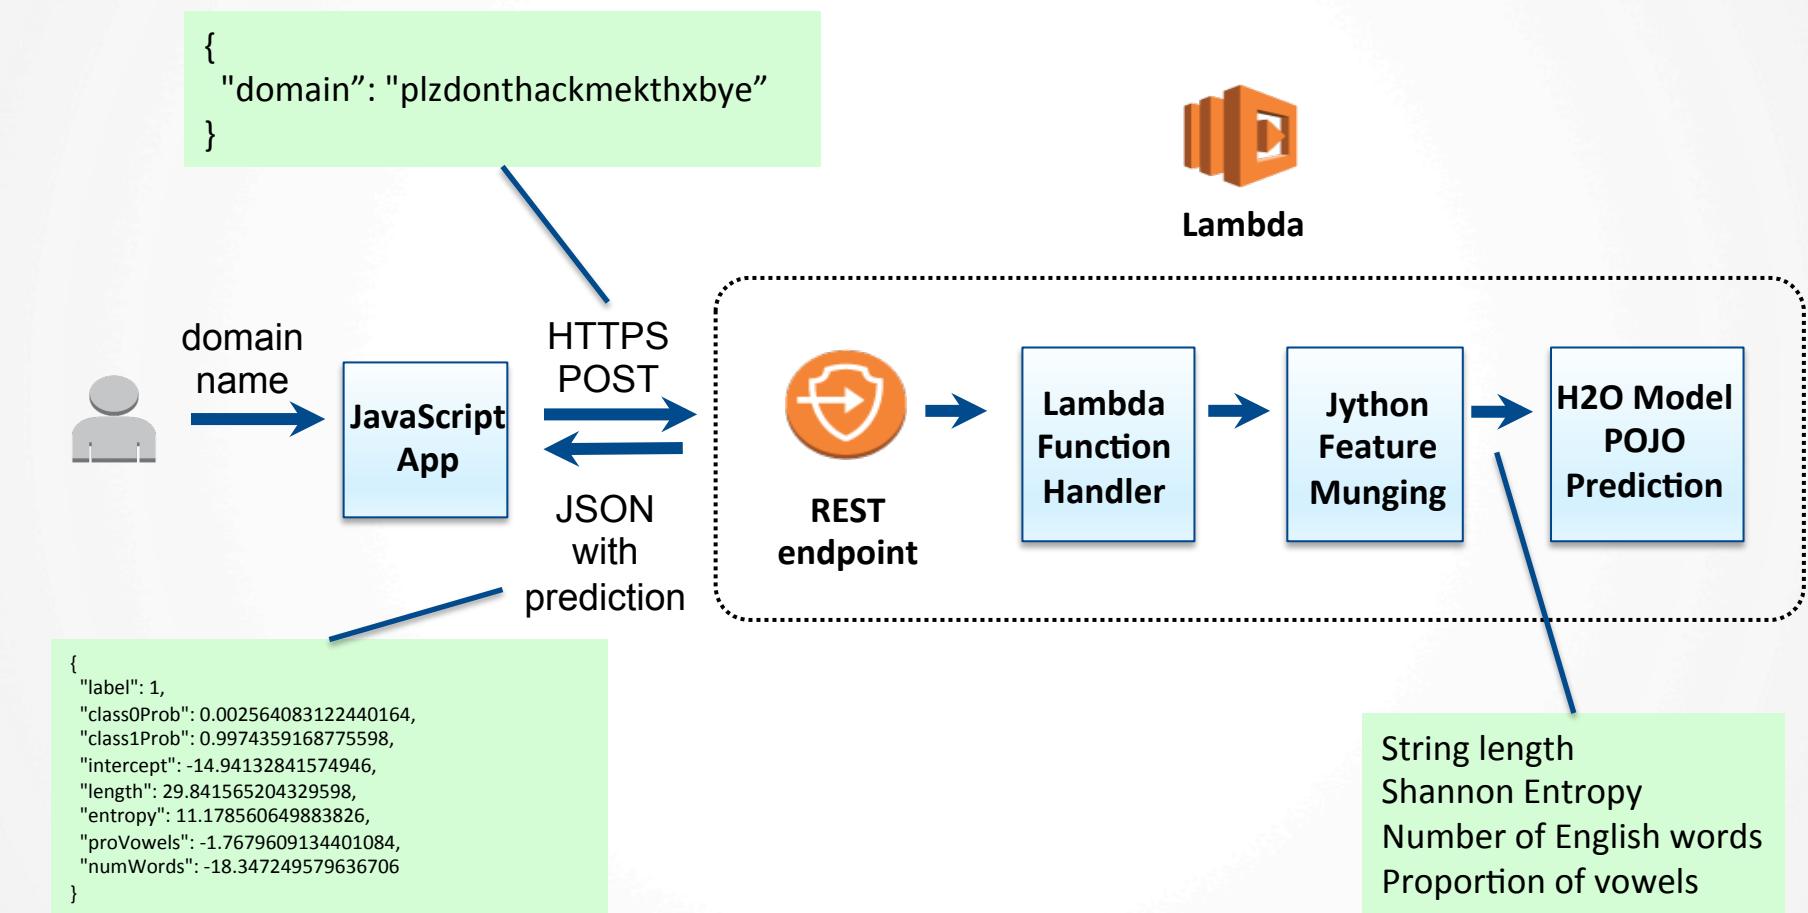
AWS LAMBDA DESIGN PATTERN

- Training data lives in data lake
- Scoring data arrives via API call
- Real-time POJO scoring in AWS Lambda

MOTIVATION FOR AWS LAMDA

USE CASE	REQUIREMENTS	TECHNOLOGY
Feature engineering	<ul style="list-style-type: none">- Use of Python for DataSci- Code re-use for training and production- Speed	Jython
Machine Learning Prediction	<ul style="list-style-type: none">- High model accuracy- Real-time environment	H2O Generated POJO Model
Deployment to Production	<ul style="list-style-type: none">- Easy handoff to Ops- Speed to deployment- Built-in scalability- No infrastructure management	AWS Lambda

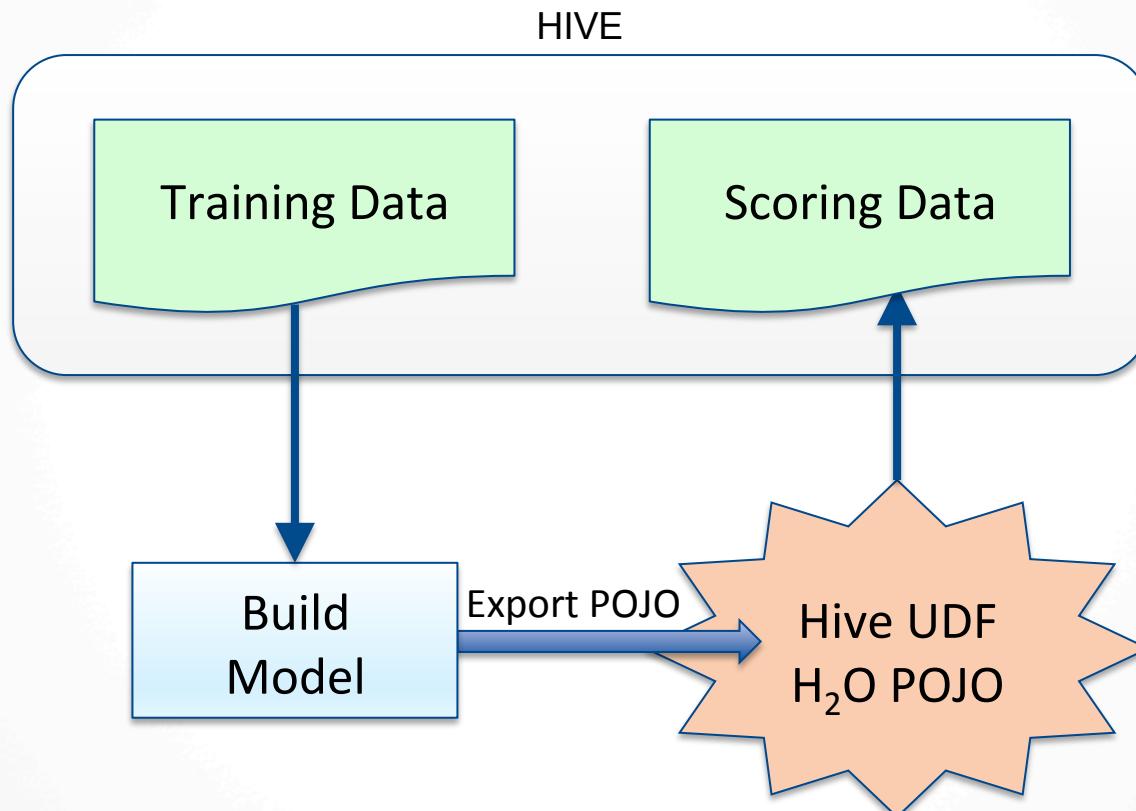
AWS LAMBDA DESIGN PATTERN (SCORING)



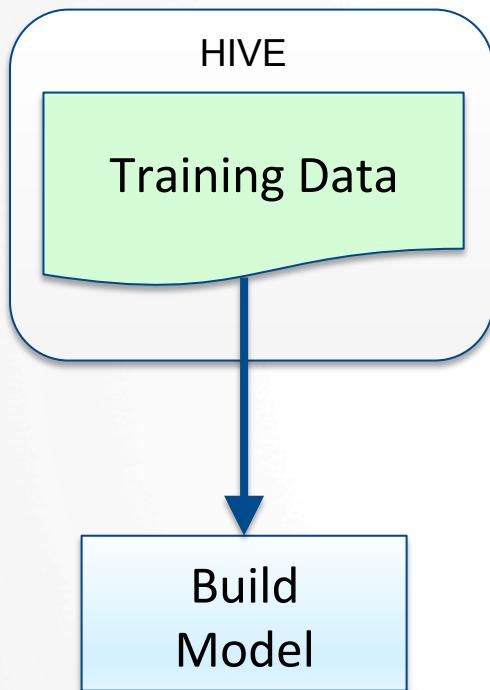
HIVE UDF DESIGN PATTERN

- Training data lives in Hive
- Scoring data lives in Hive
- Batch POJO scoring
- Models built infrequently
- Scoring occurs daily
 - Strict SLA to complete the batch job

HIVE UDF DEPLOYMENT

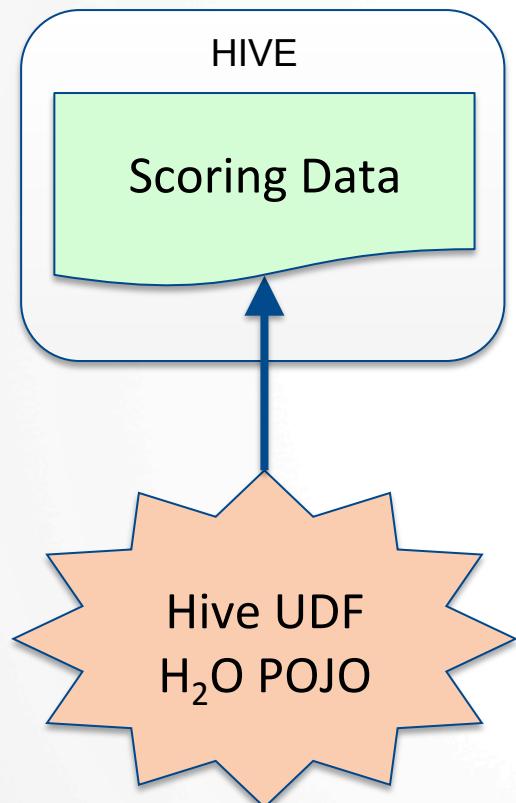


HIVE UDF (TRAINING)



- Prepare training data
 - Data may be stored in unreadable format!
 - Export to read _part files for H₂O
- Launch H₂O
- Build models
- Export models as POJOs

HIVE UDF (SCORING)



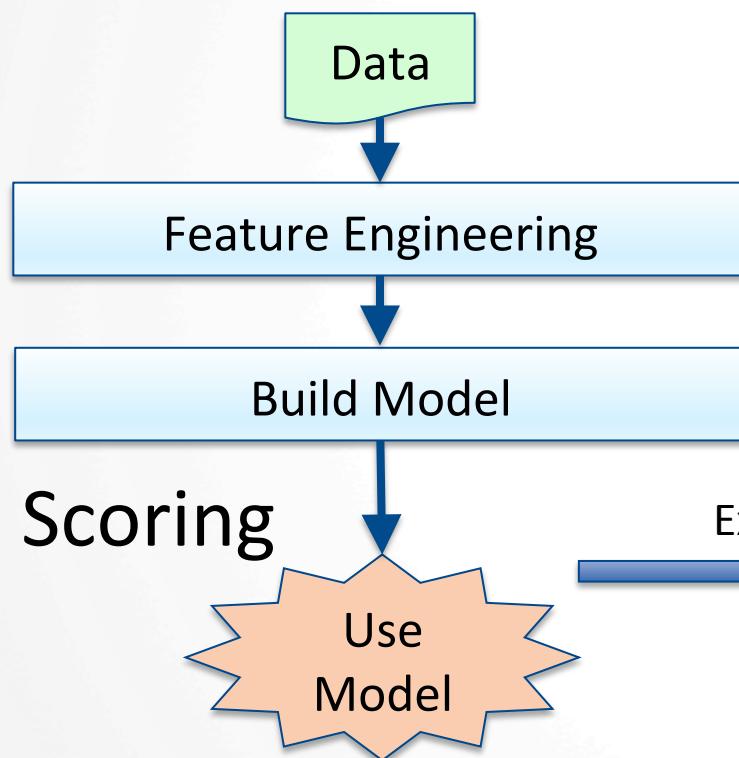
- Create Hive UDF
 - Extend GenerUDF class
 - initialize(): initialize H₂O models
 - evaluate(): loop through models & score
 - Build JAR file
- Hive
 - Load JARs, define UDF, set columns
 - ADD JAR localjars/h2o-genmodel.jar;
 - ADD JAR target/ScoreData-1.0-SNAPSHOT.jar;
 - CREATE TEMPORARY FUNCTION fn AS ...
 - set hivevar:scoredatacolnames=C1,C2...
 - Score!
 - select fn(\${scoredatacolnames}) from TABLE

JOURNALS.AI DESIGN PATTERN

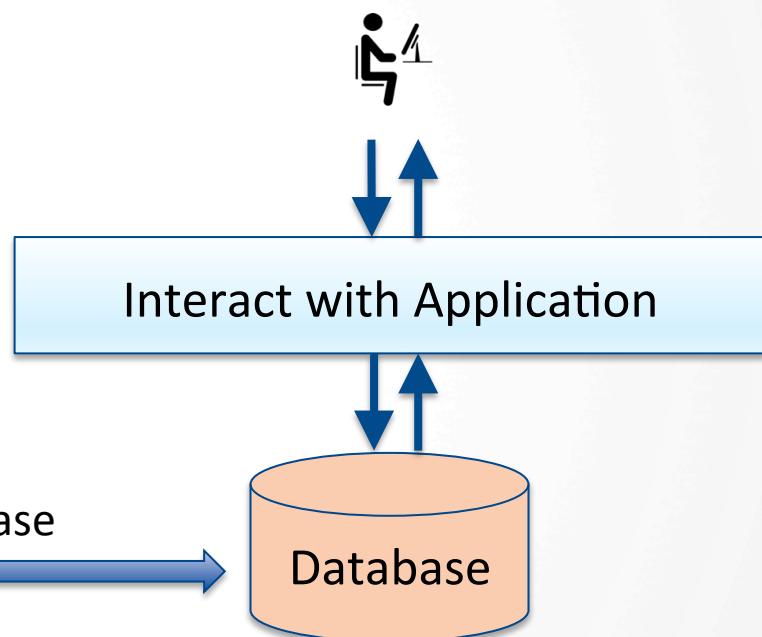
- Driven by R
- Training & scoring on bare metal servers
- Scores pre-calculated and stored in a database (e.g. MS SQL, SQLite)
- Application with domain-specific UI built on top of the database
- Flexible deployment (laptop, server, cloud, etc.)

JOURNALS.AI PRODUCTIONIZED

Training



App Deployment



DASHBOARD

TRANSACTION FLOWS

Apex, Inc.

PERIOD 01-01-2015 to 12-31-2015

Reporting currency

USD

GL system

SAP

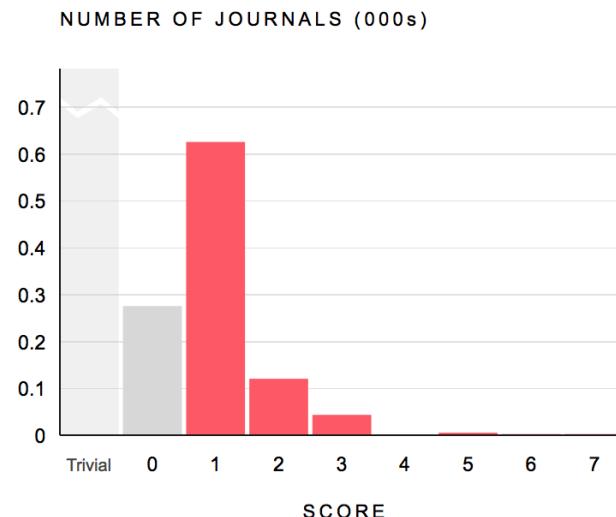
Entity type

Subsidiary of Listed Entity

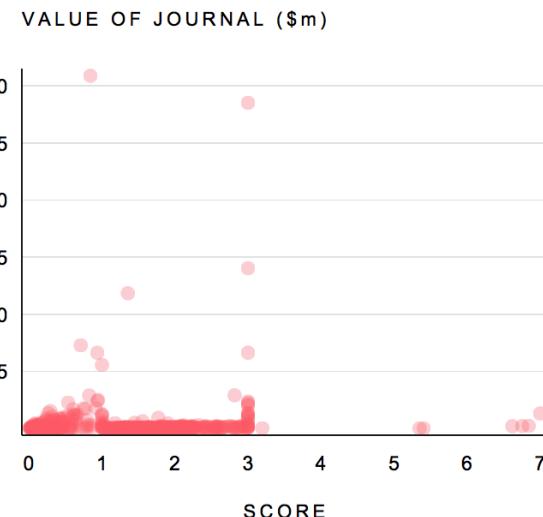
Industry

Consumer packaging

Journal by degree of unusualness



Journal by score and value



Q & A

- Thanks for attending!
- Slides in github:
 - [h2oai/h2o-meetups/tree/master/
2016_07_19_H2O_Open_Tour_NYC_Prod](https://github.com/h2oai/h2o-meetups/tree/master/2016_07_19_H2O_Open_Tour_NYC_Prod)
- Send follow up questions to:

Tom Kraljevic
tomk@h2o.ai